a. Find the titles of courses in the Comp. Sci. department that have 3 credits.

select title from course where dept name = 'Comp. Sci.' and credits = 3

b. Find the IDs of all students who were taught by an instructor named Einstein; make sure there are no duplicates in the result.

Key Idea:

- **student** connects to **takes** through <u>ID</u> (student ID).
- **instructor** connects to **teaches** through <u>ID</u> (instructor ID).
- takes and teaches connect through <u>course_id</u>, <u>sec_id</u>, <u>semester</u>, and <u>year</u>.

b. Find the IDs of all students who were taught by an instructor named Einstein; make sure there are no duplicates in the result.

Step 1: Get student-course relationships from student and takes.

? (student join takes using(ID))

Step 2: Get instructor-course relationships from instructor and teaches.

? (instructor join teaches using(ID))

Step 3: Match students and instructors through the details of course.

? (... join ... using(course_id, sec_id, semester, year))

Step 4: Filter for instructor "Einstein".

? where instructor.name = 'Einstein'

Step 5: Remove duplicates and return unique student IDs.

? select distinct student.ID

b. Find the IDs of all students who were taught by an instructor named Einstein; make sure there are <u>no</u> <u>duplicates</u> in the result.

```
select distinct student.ID

from (student join takes using(ID))

join (instructor join teaches using(ID))

using(course_id, sec_id, semester, year)

where instructor.name = 'Einstein'
```

c. Find the highest salary of any instructor.

select max(salary) from instructor

d. Find all instructors earning the highest salary (there may be more than one with the same salary).

Key Idea:

- Use an aggregate function (max) to find the highest salary in the instructor table.
- Use a **subquery** to compare each instructor's salary to the highest salary.

d. Find all instructors earning the highest salary (there may be more than one with the same salary).

- Step 1: Subquery to find the maximum salary.
- ? select max(salary) from instructor
- Step 2: Compare each instructor's salary to the maximum salary.
- ? where salary = (select max(salary) from instructor)
- Step 3: Select relevant ID and name.
- ? select ID, name from instructor

d. Find all instructors earning the highest salary (there may be more than one with the same salary).

select ID, name from instructor where salary = (select max(salary) from instructor)

e. Find the enrollment of each section that was offered in Fall 2017.

Key Idea:

- Use the section table to find all sections offered in Fall 2017.
- Use a subquery to count the number of students enrolled in each section by matching the takes table with the details of section.

e. Find the enrollment of each section that was offered in Fall 2017.

- Step 1: Filter sections offered in Fall 2017.
- ? from section where semester = 'Fall' and year = 2017
- Step 2: Select course_id and sec_id.
- ? select course_id, sec_id, ...
- Step 3: Use a subquery to count enrollment.
- (select count(ID)
 from takes
 where takes.year = section.year
 and takes.semester = section.semester
 and takes.course_id = section.course_id
 and takes.sec_id = section.sec_id)

It counts the number of students in the **takes** table who match the specific **section** with the **where** clause.

e. Find the enrollment of each section that was offered in Fall 2017.

```
select course id, sec id,
(select count(ID)
from takes
where takes.year = section.year
and takes.semester = section.semester
and takes.course id = section.course id
and takes.sec id = section.sec id)
from section where semester = 'Fall' and year = 2017
```

f. Find the maximum enrollment, across all sections, in Fall 2017.

Key Idea:

- Use a nested query to calculate the enrollment for each section.
- Use an aggregate function (max) in the outer query to find the highest enrollment.

f. Find the maximum enrollment, across all sections, in Fall 2017.

- Step 1: Filter sections offered in Fall 2017.
- ? where semester = 'Fall' and year = 2017
- Step 2: Calculate the enrollment for each section.
- select count(ID) as enrollment
 from section natural join takes
 where semester = 'Fall' and year = 2017
 group by course_id, sec_id)
- Step 3: Find the maximum enrollment.
- ? select max(enrollment) from (...)

Combines the **section** and **takes** tables based on common columns (course_id, sec_id, semester, and year) to to find the students enrolled in each section.

f. Find the maximum enrollment, across all sections, in Fall 2017.

```
select max(enrollment)
from (select count(ID) as enrollment
from section natural join takes
where semester = 'Fall' and year = 2017
group by course id, sec id)
```

g. Find the sections that had the maximum enrollment in Fall 2017.

Key Idea:

- Use a with clause to define a temporary relation.
- Use a subquery to compare each section's enrollment with the maximum enrollment and filter the results.

g. Find the sections that had the maximum enrollment in Fall 2017.

- Step 1: Define a temporary relation.
- ? with sec_enrollment as ...
- Step 2: Calculate the enrollment for each section.
- select count(ID) as enrollment
 from section natural join takes
 where semester = 'Fall' and year = 2017
 group by course_id, sec_id)
- Step 3: Find the maximum enrollment.
- ? select max(enrollment) from (...)
- Step 4: Select sections with the maximum enrollment.

select course_id, sec_id from sec_enrollment
where enrollment = (select max(enrollment) from sec_enrollment)

Combines the **section** and **takes** tables based on common columns (course_id, sec_id, semester, and year) to to find the students enrolled in each section.

g. Find the sections that had the maximum enrollment in Fall 2017.

```
with sec_enrollment as (select course_id, sec_id, count(ID) as enrollment from section natural join takes where semester = 'Fall' and year = 2017 group by course_id, sec_id) select course_id, sec_id from sec_enrollment where enrollment = (select max(enrollment) from sec_enrollment)
```

a. Increase the salary of each instructor in the Comp. Sci. department by 10%.

```
update instructor
set salary = salary * 1.10
where dept name = 'Comp. Sci.'
```

b. Delete all courses that have never been offered (i.e., do not occur in the section relation).

Key Idea:

- Use a subquery to identify courses that appear in the section table.
- Use the NOT IN operator to delete courses whose course_id does not match any course id in the section table.

b. Delete all courses that have never been offered (i.e., do not occur in the section relation).

Step 1: Find courses that have been offered.

Step 2: Find courses that have never been

offered.

Step 3: Delete courses that have never been offered.

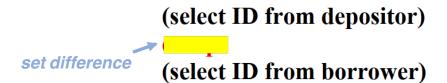
delete from course
where course_id not in
(select course_id from section)

c. Insert every student whose *tot_cred* attribute is greater than 100 as an instructor in the same department, with a salary of \$10,000.

insert into instructor select ID, name, dept_name, 10000 from student where tot cred > 100 3.8

```
branch(<u>branch_name</u>, branch city, assets)
customer (<u>ID</u>, customer_name, customer_street, customer_city)
loan (<u>loan_number</u>, branch_name, amount)
_____loan number)
account (<u>account_number</u>, branch_name, balance)
_____account_number)
```

a. Find the ID of each customer of the bank who has an account but not a loan.

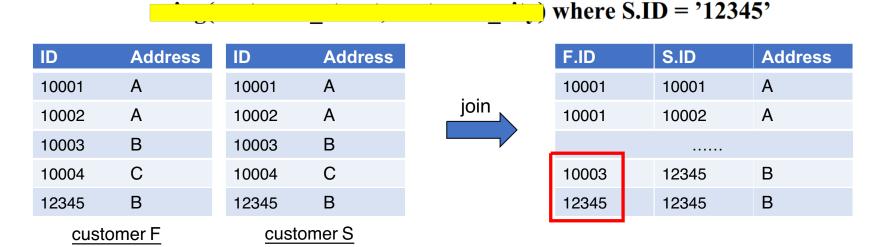


```
branch(<u>branch_name</u>, branch city, assets)

_____customer_name, ____
loan (<u>loan_number</u>, branch_name, amount)
borrower (<u>ID</u>, loan number)
account (<u>account_number</u>, branch_name, balance)
depositor (<u>ID</u>, account_number)
```

b. Find the ID of each customer who lives on the same street and in the same city as customer '12345'.

select F.ID from customer F join customer S



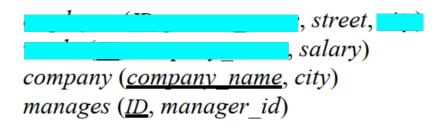
```
branch(<u>branch_name</u>, branch city, assets)

______, customer_name, customer_street, _______
loan (<u>loan_number</u>, branch_name, amount)
borrower (<u>ID</u>, loan number)

_________, balance)
```

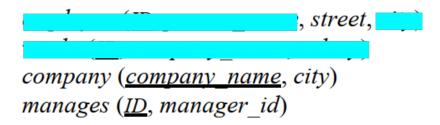
c. Find the name of each branch that has at least one customer who has an account in the bank and who lives in "Harrison".

```
select distinct branch_name from
account _______ depositor ______ customer
where customer_city = 'Harrison'
```



a. Find the ID, name, and city of residence of each employee who works for "First Bank Corporation".

select e.ID, e.person_name, e.city from employee e, works w
w.company_name = 'First Bank Corporation' and w_____



b. Find the ID, name, and city of residence of each employee who works for First Bank Corporation" and earns more than \$10000.

boolean of membership in a set

select ID, person_name, city from employee where ID (select ID from works where company name = 'First Bank Corporation' and salary > 10000)

```
, person_name, street, city)
, salary)

company (<u>company_name</u>, city)

manages (<u>ID</u>, manager_id)
```

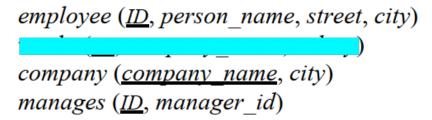
c. Find the ID of each employee who does not work for "First Bank Corporation".

select ID from works where company name <> 'First Bank Corporation'

OR

boolean of non-membership in a set

select ID from employee where ID (select ID from works where company name = 'First Bank Corporation')



d. Find the ID of each employee who earns more than every employee of "Small Bank Corporation".

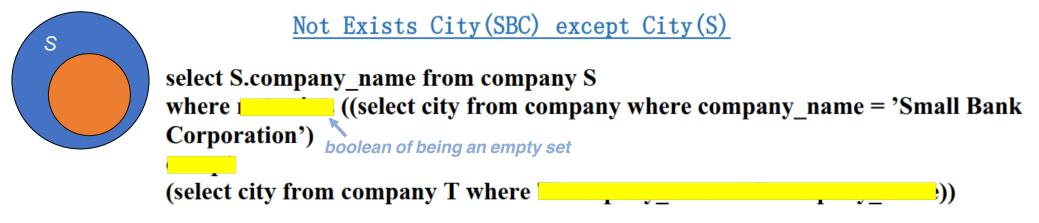
> maximum of a set

select ID from works where salary [...] (select salary from works where company name = 'Small Bank Corporation')

```
employee (<u>ID</u>, person_name, street, city)
works (<u>ID</u>, company_name, salary)

manages (<u>ID</u>, manager_id)
```

e. Assume that companies may be located in several cities. Find the name of each company that is located in every city in which "Small Bank Corporation" is located.



f. Find the name of the company that has the most employees (or companies, in the case where there is a tie for the most).



ID	company_name			
10001	A	grouping	count(distinct ID) = number of	company_name
10002	Α		employees	
10003	Α	,	3	Α
10004	В		2	В
10005	В			

g. Find the name of each company whose employees earn a higher salary, on average, than the average salary at "First Bank Corporation".

select company_name from works

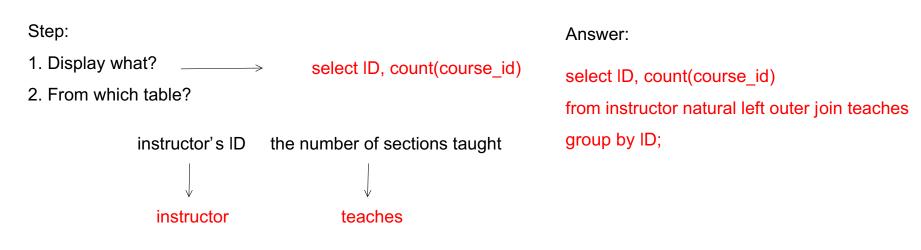
company_name	salary			
A	5000	grouping	company_name	avg(salary)
Α	6000		Α	6000
Α	7000	ŕ	В	8500
В	8000			
В	9000			

Assignment 3

(Chap 4)

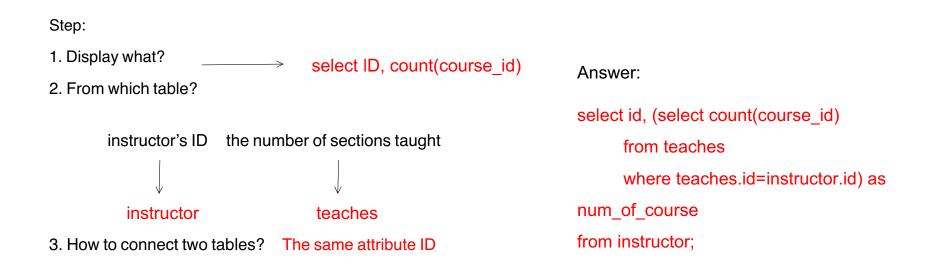
4.2 Write the following queries in SQL:

a. Display a list of all instructors, showing each instructor's ID and the number of sections taught. Make sure to show the number of sections as 0 for instructors who have not taught any section. Your query should use an outer join, and should not use subqueries.



3. The total number? group by ID

b. Write the same query as in part a, but using a scalar subquery and not using outer join.



c. Display the list of all course sections offered in Spring 2018, along with the ID and name of each instructor teaching the section. If a section has more than one instructor, that section should appear as many times in the result as it has instructors. If a section does not have any instructor, it should still appear in the result with the instructor name set to "—".

Step:

- 1. Display what? select course_id, sec_id, ID, coalesce(name, '—')
- 2. From which table? section teaches instructor
- 3. Condition? in Spring 2018

Answer:

select course_id, sec_id, ID, coalesce(name, '--')

as name

from (section natural left outer join teaches)

natural left outer join instructor

where semester = 'Spring' and year = 2018;

d. Display the list of all departments, with the total number of instructors in each department, without using subqueries. Make sure to show departments that have no instructors, and list those departments with an instructor count of zero.

Step:

- 1. Display what? select dept_name, count(ID)
- 2. From which table? department instructor
- 3. The total number? group by dept_name

Answer: select dept_name, count(ID) as instructor_num from department natural left outer join instructor

group by dept name;

4.3 Outer join expressions can be computed in SQL without using the SQL outer join operation. To illustrate this fact, show how to rewrite each of the following SQL queries without using the outer join expression.

a. select * from student natural left outer join takes

select * from student natural join takes

If we use natural join, what is the difference?

select ID, name, dept_name, tot_cred, NULL, 1.natural join: the common value on the same attribute

Answer:

union

left outer join: the common value + value only in the left table

(select ID from takes T1 where T1.id = S1.id);

NULL, NULL, NULL, NULL

from student S1 where not exists

select * from student natural join takes select from student S1 where not exists (select ID from takes T1 where T1.id = S1.id);

2. natural join: all attribute has values

left outer join: value only in the left table will show NULL on the attribution only in the right table

select ID, name, dept_name, tot_cred, NULL, NULL

b. select * from student natural full outer join takes

Full outer join = Left outer join union Right outer join

select * from student natural join takes

```
Answer:
```

```
union
(select ID, name, dept_name, tot_cred, NULL, NUL
from student S1
where not exists
(select ID from takes T1 where T1.id = S1.id))
union
(select ID, NULL, NULL, NULL, course_id, sec_id, semester, year, grade
from takes T1
where not exists
(select ID from student S1 where T1.id = S1.id));
```