



data model : method to define data structure, method to describe data.

traditional : hierarchical model
Network model

Chapter 2: Intro to Relational Model

very simple . high efficiency

Commercially used

Database System Concepts, 7th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Outline

- ① Structure of Relational Databases
- ② Database Schema
- ③ Keys
- ④ Schema Diagrams
- ⑤ Relational Query Languages
- ⑥ The Relational Algebra



Example of a *Instructor* Relation

Diagram illustrating the structure of a relation:

The table represents the *Instructor* relation with the following attributes (columns): *ID*, *name*, *dept_name*, and *salary*.

The table contains 12 tuples (rows) of data:

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

Annotations:

- attributes of the instructor*: A handwritten note pointing to the column headers.
- attributes (or columns)*: A handwritten note pointing to the column headers.
- tuples (or rows)*: A handwritten note pointing to the data rows.
- 2 dimension table*: A handwritten note describing the table as a 2-dimensional structure.



Relation Schema and Instance

- ❑ A_1, A_2, \dots, A_n are attributes
- ❑ $R = (A_1, A_2, \dots, A_n)$ is a relation schema
 - Example:
 - instructor = ($ID, name, dept_name, salary$)
schema instance
- ❑ A relation instance r defined over schema R is denoted by $r(R)$.
- ❑ The current values a relation are specified by a table
- ❑ An element t of relation r is called a tuple and is represented by a row in a table

R

ID	$name$	$dept_name$	$salary$
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000

$t \rightarrow$

r



Attributes

- ❑ The set of allowed values for each attribute is called the **domain** of the attribute
 - domain : *simple, not allowed to have complex structure*
- ❑ Attribute values are (normally) required to be **atomic**; that is, indivisible
- ❑ The special value **null** is a member of every domain. Indicated that the value is “unknown”
- ❑ The null value causes complications in the definition of many operations

table key words

```
create table instructor (
    ID          char(5),
    name        varchar(20),
    dept_name   varchar(20),
    salary      numeric(8,2))
```

↑
domain *Portuguese*



Relations are Unordered

List in Python: have order

Set: no order

- Order of tuples is irrelevant (tuples may be stored in an arbitrary order)
- Example: *instructor* relation with unordered tuples

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000



Database Schema

all the current values types

All the tables Managed by database's Schema

- ① Database schema -- is the logical structure of the database.
- ② Database instance -- is a snapshot of the data in the database at a given instant in time.
value always changes
- ③ Example:
 - schema: *instructor (ID, name, dept_name, salary)*
 - Instance: *uniquely identified*

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

ID can be
a key.

* name can be
a key.

finer (no two students
have same names)



$$R = R(A_1, A_2, \dots, A_n)$$
$$K \subseteq R$$

Keys

e.g. $A_1 \subseteq R$
 $\{A_1, A_2\} \subseteq R$
 $\{A_1, A_2, \dots, A_n\} = R \subseteq R$

- ① Let $K \subseteq R$
- ② K is a **superkey** of R if values for K are sufficient to identify a unique tuple of each possible relation $r(R)$ *e.g. {UM ID} of students is superkey*
 - Example: $\{ID\}$ and $\{ID, name\}$ are both superkeys of *instructor*.
{UM ID, name} can be reduced to {UM ID}, {UM ID} {UM ID, name} of students is also superkey.
- ③ Superkey K is a **candidate key** if K is minimal *no further reduction.* *{UM ID, name, telephone number} also superkey.*
- Example: $\{ID\}$ is a candidate key for *Instructor*
{UM ID} or {name} or {telephone}
- ④ One of the candidate keys is selected to be the **primary key**.
 - Which one? *I select.*
- ⑤ **Foreign key constraint:** Value in one relation must appear in another
 - **Referencing** relation
 - **Referenced** relation
 - Example: *dept_name* in *instructor* is a foreign key from *instructor* referencing *department*



Foreign key

- Foreign key: Value in one relation is the key in another relation
 - Referencing relation
 - Referenced relation
 - Example: *dept_name* in *instructor* is a foreign key from *instructor* referencing *department*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table **Referencing**

*values come
from another table.
foreign key values already exist there*

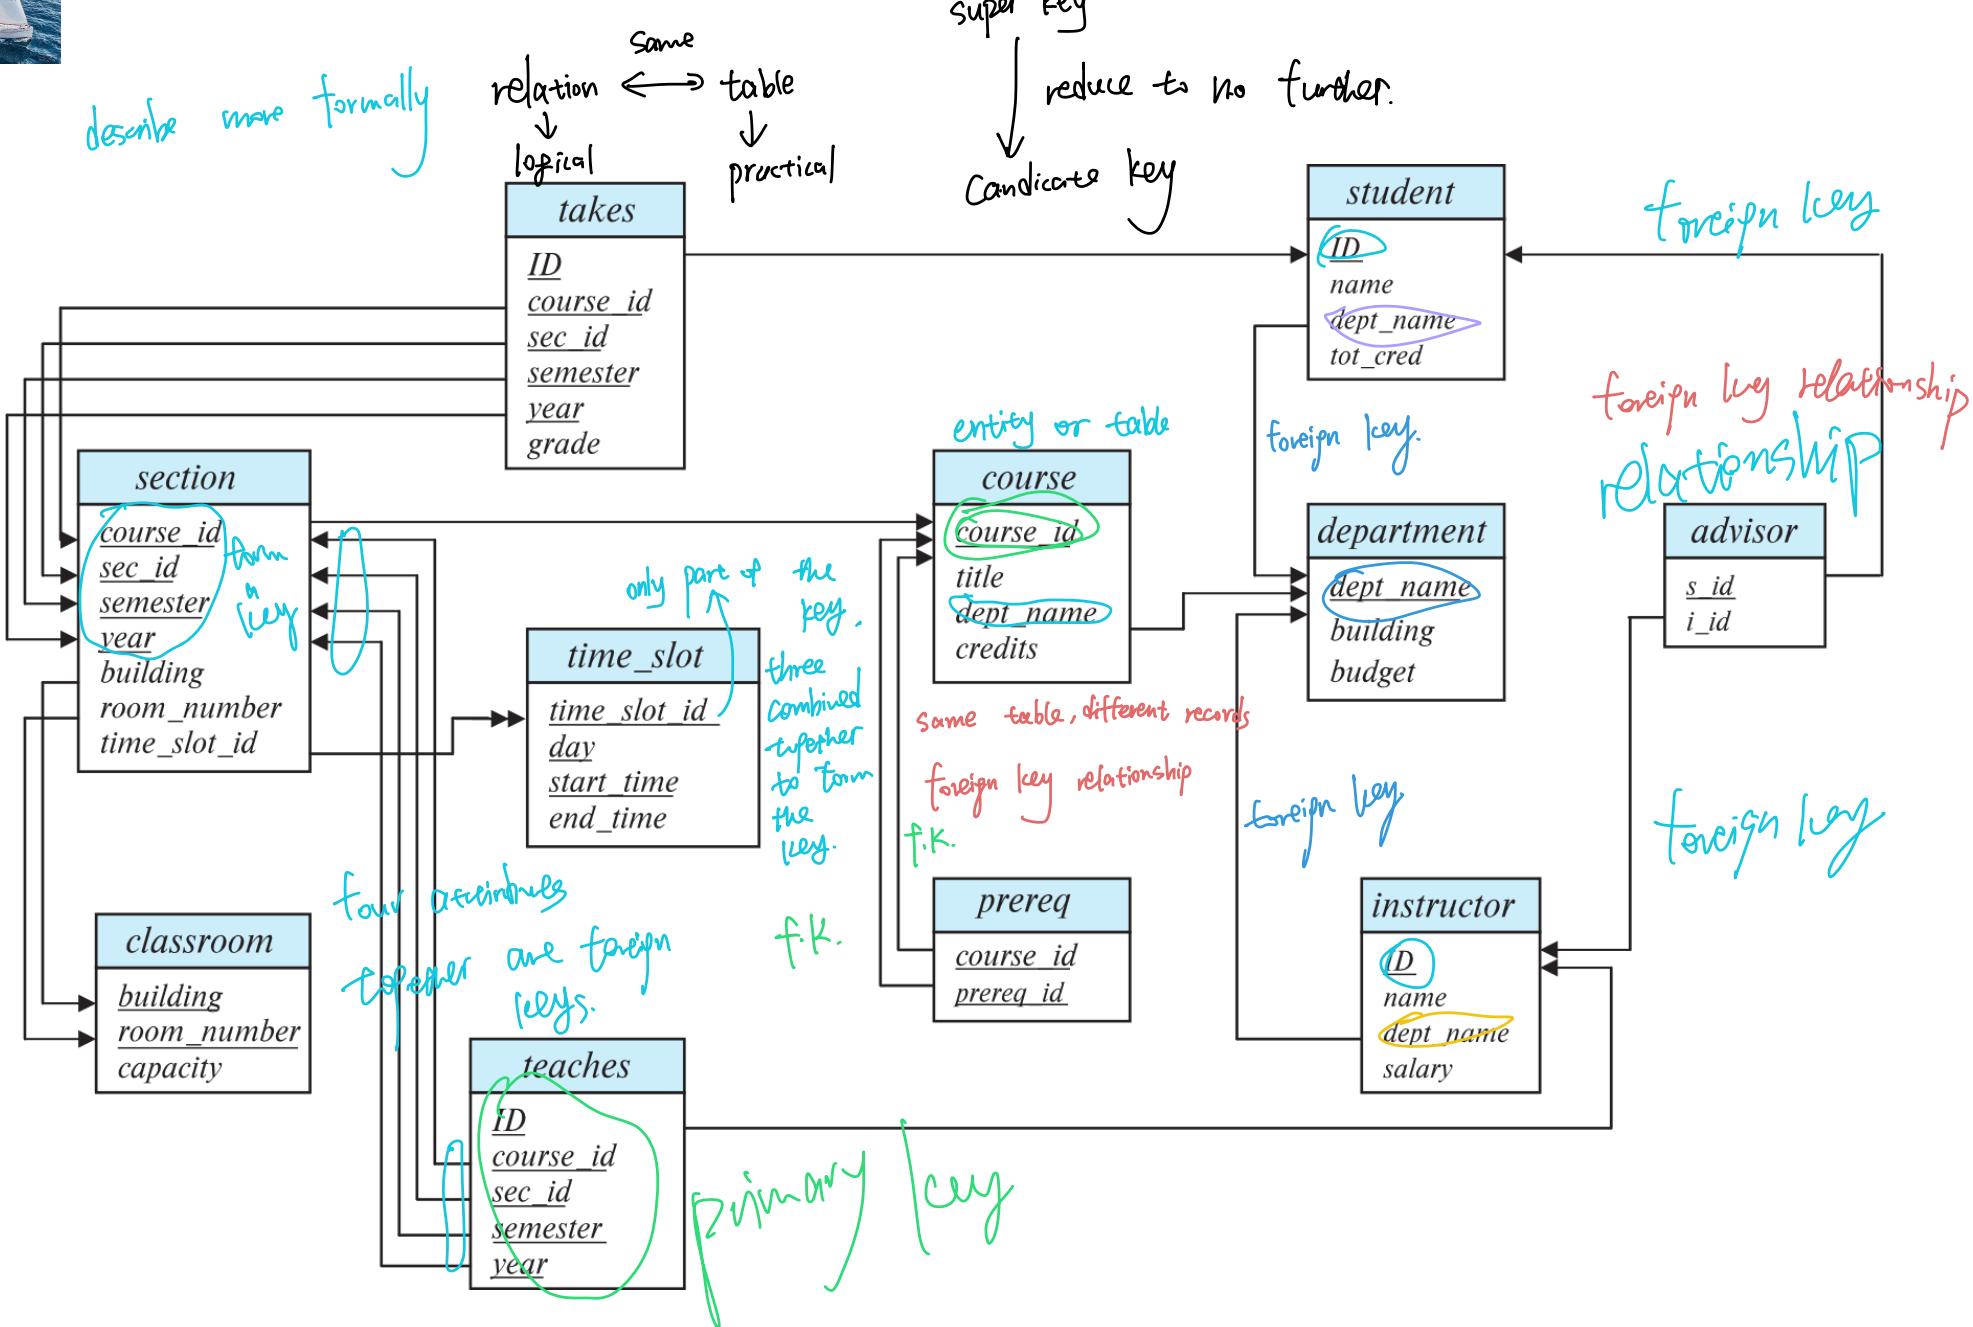
<i>dept_name</i>	<i>building</i>	<i>budget</i>
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table **Referenced**

*primary key
in this table.*



Schema Diagram for University Database





most important Relational Query Languages

Theoretical description/notation before SQL

- Procedural versus non-procedural, or declarative
- "Pure" languages:

know logical relationships

theoretical
pure
languages
for support

- The above 3 pure languages are equivalent in computing power

- We will concentrate in this chapter on relational algebra

- Not Turing-machine equivalent
- Consists of 6 basic operations

Relational algebra

- Tuple relational calculus
- Domain relational calculus

SQL some kinds of functions: equivalent, SQL practical



Relational Algebra

- ❑ A procedural language consisting of a set of operations that take one or two relations as input and produce a new relation as their result.
- ❑ Six basic operators

- select: σ Select individual record now
- project: Π only two attributes for the relation only several attributes columns
- union: \cup two sets union together
- set difference: — set 1: all the students records set 2: student records of Marcus students
Set 1 — Set 2 $S_1 = \{a, b, c, d\}$ $S_3 = S_1 - S_2 = \{c, d\}$
 $S_2 = \{a, b\}$
- Cartesian product: \times
- rename: ρ rename the attributes
 $S_4 = \{a, b, e\}$
 $S_5 = S_1 - S_4 = \{c, d\}$ because e not in S_1

Diagram illustrating the Cartesian product of relations T and S.

Relation T (t₁, t₂) has attributes A₁ and A₂. Its rows are (a, b) and (c, d).

Relation S (S₁, S₂) has attributes B₁ and B₂. Its rows are (e, f) and (x, y).

The Cartesian product T × S is shown as a single relation with attributes A₁, A₂, B₁, and B₂. It contains the following rows:

A ₁ , A ₂ , B ₁ , B ₂
a, b, e, f
a, b, x, y
c, d, e, f
c, d, x, y



Select Operation

- ? The **select** operation selects tuples that satisfy a given predicate.
- ? Notation: $\sigma_p(r)$ relation/table
- ? p is called the **selection predicate**
- ? Example: select those tuples of the *instructor* relation where the instructor is in the “Physics” department.
 - Query

$\sigma_{dept_name = "Physics"}(instructor)$
 dept-name = "Physics"

• Result
 $\sigma_{dept_name == "Physics" \text{ OR } dept_name == "CS"}(instructor)$
AND : both departments Physics & CS employ because

ID	name	dept_name	salary
22222	Einstein	Physics	95000
33456	Gold	Physics	87000

Instructor only for
one department



Select Operation (Cont.)

No " \neq " in SQL, but " $=_{<>}$ " or " $= \mid =$ "

- ④ We allow comparisons using
 $=, \neq, >, \geq, <, \leq$
in the selection predicate.
- ④ We can combine several predicates into a larger predicate by using the connectives:

④ \wedge (and) ④ \vee (or) ④ \neg (not)

- ④ Example: Find the instructors in Physics with a salary greater \$90,000, we write:

$\exists_{dept_name="Physics"} \exists_{salary > 90,000} (instructor) = \phi$ might occur

$\exists_{(dept_name='cs')}(instructor)$

- ④ The select predicate may include comparisons between two attributes.

- Example, find all departments whose name is the same as their building name:
 - $\exists_{dept_name=building} (department)$



Project Operation

- ❑ A unary operation that returns its argument relation, with certain attributes left out.
- ❑ Notation:

$$\prod_{A_1, A_2, A_3 \dots A_k} (r)$$

where A_1, A_2, \dots, A_k are attribute names and r is a relation name.

- ❑ The result is defined as the relation of k columns obtained by erasing the columns that are not listed
- ❑ Duplicate rows removed from result, since relations are sets

↓

ID	name	GPA
1111	Eric	3.0
2222	Eric	3.0

Project

name	GPA
Eric	3.0
Eric	3.0

result is a form of set
conceptual
clear
remove duplicate



Project Operation Example

- Example: eliminate the *dept_name* attribute of *instructor*
- Query:

$\pi_{ID, name, salary} (instructor)$

- Result:

<i>ID</i>	<i>name</i>	<i>salary</i>
10101	Srinivasan	65000
12121	Wu	90000
15151	Mozart	40000
22222	Einstein	95000
32343	El Said	60000
33456	Gold	87000
45565	Katz	75000
58583	Califieri	62000
76543	Singh	80000
76766	Crick	72000
83821	Brandt	92000
98345	Kim	80000



Composition of Relational Operations

- ② The result of a relational-algebra operation is relation and therefore of relational-algebra operations can be composed together into a relational-algebra expression.
- ② Consider the query -- Find the names of all instructors in the Physics department.

$\prod_{\text{dept_name} = \text{"Physics"} \text{ (instructor)}} \text{name}$

procedure
cannot change order
if we name first, $\text{dept} = \text{physics}$

- ② Instead of giving the name of a relation as the argument of the projection operation, we give an expression that evaluates to a relation.

can't
be done



Cartesian-Product Operation

- ② The Cartesian-product operation (denoted by X) allows us to combine information from any two relations.
- ② Example: the Cartesian product of the relations *instructor* and *teaches* is written as:

$$\textit{instructor} \times \textit{teaches} \quad |_{\text{100}} \times |_{\text{200}} = |_{\text{20000}}$$

- ② We construct a tuple of the result out of each possible pair of tuples: one from the *instructor* relation and one from the *teaches* relation (see next slide)
- ② Since the instructor *ID* appears in both relations we distinguish between these attribute by attaching to the attribute the name of the relation from which the attribute originally came.

- *instructor.ID* |¹⁰⁰
- *teaches.ID* |²⁰⁰



The *instructor* X *teaches* table

Select: $\text{G}_{\text{instructor.ID} = \text{teaches.ID}}(\text{IXT})$

join:

$\text{I} \bowtie_{\text{id}} \text{T}$

<i>instructor.ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>teaches.ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2017
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2017
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2017
...
...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2017
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2018
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2017
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2018
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2018
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2017
...
...
15151	Mozart	Music	40000	10101	CS-101	1	Fall	2017
15151	Mozart	Music	40000	10101	CS-315	1	Spring	2018
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2017
15151	Mozart	Music	40000	12121	FIN-201	1	Spring	2018
15151	Mozart	Music	40000	15151	MU-199	1	Spring	2018
15151	Mozart	Music	40000	22222	PHY-101	1	Fall	2017
...
...
22222	Einstein	Physics	95000	10101	CS-101	1	Fall	2017
22222	Einstein	Physics	95000	10101	CS-315	1	Spring	2018
22222	Einstein	Physics	95000	10101	CS-347	1	Fall	2017
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2018
22222	Einstein	Physics	95000	15151	MU-199	1	Spring	2018
22222	Einstein	Physics	95000	22222	PHY-101	1	Fall	2017
...
...



Join Operation

❑ The Cartesian-Product

instructor X teaches

associates every tuple of *instructor* with every tuple of *teaches*.

- Most of the resulting rows have information about instructors who did NOT teach a particular course.

❑ To get only those tuples of “*instructor X teaches*” that pertain to instructors and the courses that they taught, we write:

$\delta_{instructor.id = teaches.id} (instructor \times teaches)$

- We get only those tuples of “*instructor X teaches*” that pertain to instructors and the courses that they taught.

❑ The result of this expression, shown in the next slide



Join Operation (Cont.)

- ② The table corresponding to:

$$\text{σ}_{instructor.id = teaches.id} (instructor \times teaches)$$

<i>instructor.ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>teaches.ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2017
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2018
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2017
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2018
15151	Mozart	Music	40000	15151	MU-199	1	Spring	2018
22222	Einstein	Physics	95000	22222	PHY-101	1	Fall	2017
32343	El Said	History	60000	32343	HIS-351	1	Spring	2018
45565	Katz	Comp. Sci.	75000	45565	CS-101	1	Spring	2018
45565	Katz	Comp. Sci.	75000	45565	CS-319	1	Spring	2018
76766	Crick	Biology	72000	76766	BIO-101	1	Summer	2017
76766	Crick	Biology	72000	76766	BIO-301	1	Summer	2018
83821	Brandt	Comp. Sci.	92000	83821	CS-190	1	Spring	2017
83821	Brandt	Comp. Sci.	92000	83821	CS-190	2	Spring	2017
83821	Brandt	Comp. Sci.	92000	83821	CS-319	2	Spring	2018
98345	Kim	Elec. Eng.	80000	98345	EE-181	1	Spring	2017



Join Operation (Cont.)

- ? The **join** operation allows us to combine a select operation and a Cartesian-Product operation into a single operation.
- ? Consider relations $r (R)$ and $s (S)$
- ? Let “theta” be a predicate on attributes in the schema R “union” S . The join operation $r \theta s$ is defined as follows:

$$r \theta s = \sigma_{\theta}(r \times s)$$

- ? Thus

$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$

- ? Can equivalently be written as

~~instructor~~ $_{Instructor.id = teaches.id} \sigma_{Instructor.id = teaches.id} teaches.$
more generalized

instructor $_{Instructor.id = teaches.id} \sigma_{Instructor.id = teaches.id} teachers.$

or instructor ... > ... teachers.



Union Operation

- ❑ The union operation allows us to combine two relations
- ❑ Notation: $r \cup s$
- ❑ For $r \cup s$ to be valid.
 1. r, s must have the same **arity** (same number of attributes)
 2. The attribute domains must be **compatible** (example: 2nd column of r deals with the same type of values as does the 2nd column of s)
- ❑ Example: to find all courses taught in the Fall 2017 semester, or in the Spring 2018 semester, or in both

$$\exists course_id (\exists semester = "Fall" \wedge year = 2017 (section)) \cup \exists course_id (\exists semester = "Spring" \wedge year = 2018 (section))$$

Schema incompatible
different



Union Operation (Cont.)

Result of:

$$\begin{array}{l} \text{U}_{course_id} (\text{G}_{semester="Fall"} \wedge year=2017 (section)) \cup \\ \text{U}_{course_id} (\text{G}_{semester="Spring"} \wedge year=2018 (section)) \end{array}$$

course_id
CS-101
CS-315
CS-319
CS-347
FIN-201
HIS-351
MU-199
PHY-101



Set-Intersection Operation

- ❑ The set-intersection operation allows us to find tuples that are in both the input relations.
- ❑ Notation: $r \cap s$
- ❑ Assume:
 - r, s have the same arity
 - attributes of r and s are compatible
- ❑ Example: Find the set of all courses taught in both the Fall 2017 and the Spring 2018 semesters.

$\exists course_id (\exists semester = "Fall" \wedge year = 2017 (section)) \cap \exists course_id (\exists semester = "Spring" \wedge year = 2018 (section))$

- Result

course_id
CS-101



Set Difference Operation

- ❑ The set-difference operation allows us to find tuples that are in one relation but are not in another.
- ❑ Notation $r - s$ *in r but not in s.*
- ❑ Set differences must be taken between **compatible** relations.
 - r and s must have the same arity
 - attribute domains of r and s must be compatible
- ❑ Example: to find all courses taught in the Fall 2017 semester, but not in the Spring 2018 semester

$\exists \text{course_id} (\text{semester} = \text{"Fall"} \wedge \text{year} = 2017 \text{ (section)}) -$
 $\exists \text{course_id} (\text{semester} = \text{"Spring"} \wedge \text{year} = 2018 \text{ (section)})$

for each element in A,
check whether in B and
remains if not in B.

course_id
CS-347
PHY-101

*CS-101 both in A and B
if A - B,
CS-101 still eliminated.*



The Assignment Operation

- ❑ It is convenient at times to write a relational-algebra expression by assigning parts of it to temporary relation variables.
- ❑ The assignment operation is denoted by \Leftarrow and works like assignment in a programming language.
- ❑ Example: Find all instructor in the “Physics” and Music department.

Physics \Leftarrow $\{ \text{dept_name} = \text{“Physics”} \}$ (*instructor*)

Music \Leftarrow $\{ \text{dept_name} = \text{“Music”} \}$ (*instructor*)

Physics \sqcup *Music*

- ❑ With the assignment operation, a query can be written as a sequential program consisting of a series of assignments followed by an expression whose value is displayed as the result of the query.



The Rename Operation

- ② The results of relational-algebra expressions do not have a name that we can use to refer to them. The rename operator, ρ is provided for that purpose

- ② The expression:

$$\rho_x(E)$$

returns the result of expression E under the name x

- ② Another form of the rename operation:

$$\rho_{x(A_1, A_2, \dots, A_n)}(E)$$

↓ attributes' new name

↓ table's new name

if P name,
instructor is seen the
name of the transformed
tables, e.g.,

$$R = \text{Instructor}(ID, name, Dept-name-Salary)$$
$$\rho_{Ins [ID, name, Dept-name, salary]}(R)$$



Equivalent Queries

- ❑ There is more than one way to write a query in relational algebra.
- ❑ Example: Find information about courses taught by instructors in the Physics department with salary greater than 90,000
- ❑ Query 1

$\exists \text{dept_name} = "Physics" \forall \text{salary} > 90,000 \text{ (instructor)}$ better in terms of complexity

- ❑ Query 2

$\exists \text{dept_name} = "Physics" (\exists \text{salary} > 90,000 \text{ (instructor)})$ main memory, second memory enter main memory and hard disk twice.

- ❑ The two queries are not identical; they are, however, equivalent -- they give the same result on any database.

DBMS: for 2 queries, will choose the first one



Equivalent Queries

Not User But DBMS will evaluate which query is computationally better.

- ❑ There is more than one way to write a query in relational algebra.
- ❑ Example: Find information about courses taught by instructors in the Physics department
- ❑ Query 1 *big table as join is very expensive*
~~($\sigma_{dept_name='Physics'}$ (instructor) \bowtie instructor.ID = teaches.ID teaches)~~
- ❑ Query 2 *better first reduce no. of rows of the large table.*
~~($\sigma_{dept_name='Physics'}$ (instructor)) \bowtie instructor.ID = teaches.ID teaches~~
- ❑ The two queries are not identical; they are, however, equivalent -- they give the same result on any database.



End of Chapter 2