

## Set Membership

Database System Concepts – 7th Edition

3.38

©SBerchertz, North and Sudarshan

## Set Comparison

Database System Concepts – 7th Edition

3.39

©SBerchertz, North and Sudarshan

## Set Comparison – “all” Clause

- ¼ Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

```
select name
  from instructor
 where salary > all (select salary
                    from instructor
                    where dept_name = 'Biology');
```

Database System Concepts – 7th Edition

3.42

©SBerchertz, North and Sudarshan

## Set Membership

- ¼ Find courses offered in Fall 2017
- ```
select distinct course_id
  from section
 where semester = 'Fall' and year = 2017;
```
- ¼ Find courses offered in Fall 2017 and in Spring 2018
- ```
select distinct course_id
  from section
 where semester = 'Fall' and year = 2017 and
        course_id in (select course_id
                      from section
                      where semester = 'Spring' and year = 2018);
```
- ¼ Find courses offered in Fall 2017 but not in Spring 2018
- ```
select distinct course_id
  from section
 where semester = 'Fall' and year = 2017 and
        course_id not in (select course_id
                          from section
                          where semester = 'Spring' and year = 2018);
```

Database System Concepts – 7th Edition

3.37

©SBerchertz, North and Sudarshan

## Set Comparison – “some” Clause

- ¼ Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name
  from instructor as T, instructor as S
 where T.salary > S.salary and S.dept_name = 'Biology';
```

- ¼ Same query using > some clause

```
select name
  from instructor
 where salary > some (select salary
                     from instructor
                     where dept_name = 'Biology');
```

Database System Concepts – 7th Edition

3.40

©SBerchertz, North and Sudarshan

## Definition of “all” Clause

- ¼  $F \text{ <comp> all } r_1 \text{ } t \text{ } r_2 \text{ (} F \text{ <comp> } t \text{)}$

$(5 < \text{all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$

$(5 < \text{all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$

$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 \geq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \geq 4 \text{ and } 5 \geq 6)$

$(\forall \text{all}) \text{ } \forall \text{not in}$   
However,  $(= \text{all}) \text{ } \nexists \text{ in}$

Database System Concepts – 7th Edition

3.43

©SBerchertz, North and Sudarshan

## Set Membership (Cont.)

- ¼ Name all instructors whose name is neither “Mozart” nor Einstein”
- ```
select distinct name
  from instructor
 where name not in ('Mozart', 'Einstein')
```
- ¼ Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101
- ```
select count (distinct ID)
  from takes
 where (course_id, sec_id, semester, year) in
        (select course_id, sec_id, semester, year
         from teaches
         where teaches.ID = 10101);
```
- ¼ Note: Above query can be written in a much simpler manner. The formulation above is simply to illustrate SQL features

Database System Concepts – 7th Edition

3.38

©SBerchertz, North and Sudarshan

## Definition of “some” Clause

- ¼  $F \text{ <comp> some } r_1 \text{ } t \text{ } r_2 \text{ such that (} F \text{ <comp> } t \text{)}$   
Where <comp> can be:

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$  (read:  $5 < \text{some tuple in the relation}$ )

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \geq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \geq 5)$

$(= \text{some}) \text{ } \forall \text{in}$   
However,  $(\exists \text{some}) \text{ } \nexists \text{ in}$

Database System Concepts – 7th Edition

3.41

©SBerchertz, North and Sudarshan

## Test for Empty Relations

- ¼ The exists construct returns the value true if the argument subquery is nonempty.
- ¼ exists  $r \text{ } \forall r \neq \emptyset$
- ¼ not exists  $r \text{ } \forall r = \emptyset$

Database System Concepts – 7th Edition

3.44

©SBerchertz, North and Sudarshan

### Use of "exists" Clause

- ¼ Yet another way of specifying the query "Find all courses taught in both the Fall 2017 semester and in the Spring 2018 semester"
- ```
select course_id
from section as S
where semester = 'Fall' and year = 2017 and
exists (select *
        from section as T
        where semester = 'Spring' and year= 2018
        and S.course_id = T.course_id);
```
- ¼ Correlation name – variable S in the outer query
  - ¼ Correlated subquery – the inner query
  - ¼ "Find all courses taught in the Fall 2017 semester but not in the Spring 2018 semester"
- How to revise above query?
  - exist ¼ not exist

Database System Concepts – 7th Edition

3.45

©SBShirshat, North and Subrahman

### Subqueries in the From Clause

### Complex Queries using With Clause

- ¼ Find all departments where the total salary is greater than the average of the total salary at all departments
- ```
with dept_total (dept_name, value) as
(select dept_name, sum(salary)
 from instructor
 group by dept_name),
dept_total_avg (value) as
(select avg(value)
 from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value > dept_total_avg.value;
```
- As many as temporary table in with clause could be excluded*
- only one main body*

Database System Concepts – 7th Edition

3.51

©SBShirshat, North and Subrahman

### Use of "not exists" Clause

- ¼ Find all students who have taken all courses offered in the Biology department.
- ```
select distinct S.ID, S.name
from student as S
where not exists ((select course_id
                  from course
                  where dept_name = 'Biology')
except
(select T.course_id
 from takes as T
 where S.ID = T.ID));
```
- First nested query lists all courses offered in Biology
  - Second nested query lists all courses a particular student took
- ¼ Note that  $X - Y = \emptyset$  ¼  $X \subseteq Y$
  - ¼ Note: Cannot write this query using = all and its variants

Database System Concepts – 7th Edition

3.46

©SBShirshat, North and Subrahman

### Subqueries in the From Clause

- ¼ SQL allows a subquery expression to be used in the from clause
  - ¼ Find the average instructors' salaries of those departments where the average salary is greater than \$42,000:
- ```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
      from instructor
      group by dept_name) as dept_avg
where avg_salary > 42000;
```
- ¼ Note that we do not need to use the having clause
  - ¼ Another way to write above query
- ```
select dept_name, avg_salary
from (select dept_name, avg (salary)
      from instructor
      group by dept_name)
as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```

Database System Concepts – 7th Edition

3.49

©SBShirshat, North and Subrahman

### Scalar Subquery

- ¼ Scalar subquery is one which is used where a single value is expected
  - ¼ List the number of instructors in 'CIS' department
- ```
select count(*)
from instructor
where instructor.dept_name='Com.Sci';
```
- ¼ List all departments along with the number of instructors in each department
- ```
select dept_name,
(select count(*)
 from instructor
 where instructor.dept_name= department.dept_name )
as num_instructors
from department;
```
- main body*
- query returns to find /k/ of instructors in each department.*
- ¼ Runtime error if subquery returns more than one result tuple

Database System Concepts – 7th Edition

3.52

©SBShirshat, North and Subrahman

### Test for Absence of Duplicate Tuples

- ¼ The unique construct tests whether a subquery has any duplicate tuples in its result.
  - ¼ The unique construct evaluates to "true" if a given subquery contains no duplicates.
  - ¼ Find all courses that were offered at most once in 2017
- ```
select T.course_id
from course as T
where unique (select R.course_id
             from section as R
             where T.course_id= R.course_id
             and R.year = 2017);
```
- ¼ Remark: MySQL not support unique query

Database System Concepts – 7th Edition

3.47

©SBShirshat, North and Subrahman

### With Clause

- ¼ The with clause provides a way of defining a temporary relation whose definition is available only to the query in which the with clause occurs.
  - ¼ Find all departments with the maximum budget
- ```
with max_budget (dept_name, value) as
(select max(budget)
 from department)
select department.dept_name
from department, max_budget
where department.budget = max_budget.value;
```
- max\_budget (value) is a temporary table*
- value is a scalar value*

Database System Concepts – 7th Edition

3.50

©SBShirshat, North and Subrahman

### Modification of the Database

- ¼ Deletion of tuples from a given relation.
- ¼ Insertion of new tuples into a given relation
- ¼ Updating of values in some tuples in a given relation

Database System Concepts – 7th Edition

3.53

©SBShirshat, North and Subrahman

## Deletion

- ¾ Delete all instructors  
delete from instructor;
- ¾ Delete all instructors from the Finance department  
delete from instructor  
where dept\_name= 'Finance';
- ¾ Delete all tuples in the instructor relation for those instructors associated with a department located in the Watson building.  
delete from instructor  
where dept\_name in (select dept\_name  
from department  
where building = 'Watson');

Database System Concepts - 7th Edition

3.54

©SBBaratz, North and Subrahman

## Insertion (Cont.)

- ¾ Make each student in the Music department who has earned more than 144 credit hours an instructor in the Music department with a salary of \$18,000.  
insert into instructor *addition to instructor*  
select ID, name, dept\_name, 18000 *data tuple student*  
from student  
where dept\_name = 'Music' and total\_cred > 144;
- ¾ The select from where statement is evaluated fully before any of its *conds are inserted into the relation*  
Otherwise queries like  
insert into table1 select \* from table1  
would cause problem

Database System Concepts - 7th Edition

3.57

©SBBaratz, North and Subrahman

## Case Statement for Conditional Updates

- ¾ Same query as before but with case statement  
update instructor *if case in then*  
set salary = *case* *when salary <= 100000 then salary \* 1.05*  
else salary \* 1.03  
end

Database System Concepts - 7th Edition

3.60

©SBBaratz, North and Subrahman

## Deletion (Cont.)

- ¾ Delete all instructors whose salary is less than the average salary of instructors  
delete from instructor  
where salary < (select avg (salary)  
from instructor);
- Problem: as we delete tuples from instructor, the average salary changes
- Solution used in SQL:
  1. First, compute avg (salary) and find all tuples to delete
  2. Next, delete all tuples found above (without recomputing avg or retesting the tuples)

Database System Concepts - 7th Edition

3.55

©SBBaratz, North and Subrahman

## Updates

- ¾ Give a 5% salary raise to all instructors  
update instructor  
set salary = salary \* 1.05
- ¾ Give a 5% salary raise to those instructors who earn less than 70000  
update instructor  
set salary = salary \* 1.05  
where salary < 70000;
- ¾ Give a 5% salary raise to instructors whose salary is less than average  
update instructor  
set salary = salary \* 1.05  
where salary < (select avg (salary)  
from instructor);

Database System Concepts - 7th Edition

3.58

©SBBaratz, North and Subrahman

## Updates with Scalar Subqueries

- ¾ Recompute and update tot\_creds value for all students  
update student S  
set tot\_cred = (select sum(credits)  
from takes, course  
where takes.course\_id = course.course\_id and  
S.ID= takes.ID and  
takes.grade <> 'F' and  
takes.grade is not null);
- ¾ Sets tot\_creds to null for students who have not taken any course
- ¾ Instead of sum(credits), use:  
select case  
when sum(credits) is not null then sum(credits)  
else 0  
end

Database System Concepts - 7th Edition

3.61

©SBBaratz, North and Subrahman

## Insertion

- ¾ Add a new tuple to course  
insert into course *data tuple for this insert*  
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4),  
('CS-456', 'Data Mining', 'Comp. Sci.', 4);
- ¾ or equivalently  
insert into course (course\_id, title, dept\_name, credits)  
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4),  
('CS-456', 'Data Mining', 'Comp. Sci.', 4);
- ¾ Add a new tuple to student with tot\_creds set to null  
insert into student  
values ('3003', 'Green', 'Finance', null);

Database System Concepts - 7th Edition

3.56

©SBBaratz, North and Subrahman

## Updates (Cont.)

- ¾ Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
  - Write two update statements:  
update instructor  
set salary = salary \* 1.03  
where salary > 100000;  
update instructor  
set salary = salary \* 1.05  
where salary <= 100000;
  - The order is important
  - Can be done better using the case statement (next slide)

Database System Concepts - 7th Edition

3.59

©SBBaratz, North and Subrahman

End of Chapter 3

Database System Concepts - 7th Edition

3.62

©SBBaratz, North and Subrahman